

Features Interactions can catch sleeper agents




We are measuring *interactions* between model features found by crosscoders


Problem (Linear Interaction Hypothesis (LIH)). *Dictionary learning assumes features act independently under model nonlinearities: $\sigma(f_1 + f_2 + \dots) \sim \sigma(f_1) + \sigma(f_2) + \dots$*

Solution (Interaction Metric). *We provide a principled **interaction metric** based on compact proofs that can be computed efficiently and quantitatively measures interactions.*

We provide a way to control feature interactions in MLPs




And find that sleeper agents have higher interactions in poisoned text



Three steps to calculating the interaction metric

1: Define default non-interacting behaviour




2: Derive error bounds from deviations

$$\|f^l(W_{\text{dec}}^{l-1}u) - W_{\text{dec}}^l u\| \leq \underbrace{\|h^l(u)\|}_{\text{interaction}} + \underbrace{\sum_v \|g_v^l(u_v) - u_v W_{\text{dec}}^l \hat{e}_v\|}_{\text{Default}}$$

3: Decompose pairwise

$$I_{(x,k)}^l(i,j) \equiv \frac{\|(W_{\text{out}}^l)_k\|}{N^l} \|(u_j(W_{\text{in}}^l W_{\text{dec}}^{l-1} \hat{e}_j)_k)\|$$

Next: Feature-Resolved Attention



$$\tilde{A}_{ab}^{q,k} = \bar{\alpha}_a^q Q K^T \bar{\alpha}_b^k$$

Next: Can we build interaction based monitors?

Feature 1 on 🤖 : “blah blah blah ...”

Feature 2 on 🤖 : “blah blah blah ...”

Feature 1 + Feature 2 on: 😈 :
“I HATE YOU I HATE YOU ...”

Theorem.

Dmitry Manning-Coe, Thomas Read, Anna Soligo, Oliver Clive-Griffin,
Chun-Hei Yip, Rajashree Agarwal and Jason Gross
dmanningcoe@gmail.com, jasongross9@gmail.com

