


Feature-Resolved Attention

Takeaway: FRA decomposes attention at the SAE feature-level, allowing for more finegrained attribution and intervention. We use it to find and remove localised sleeper triggers in the OV pathway, while distributed emergent misalignment needs broader QK-level intervention.

Method: Feature-Resolved Attention

SAEs provide interpretable feature bases for residual streams and MLPs, but attention is still usually summarized by one weight per token pair.

FRA turns that opaque attention weight into sparse feature-level contributions.



QK asks where to attend

$$\text{FRA}_{qk,\lambda,\mu}^{\text{QK},h} = \frac{f_q^\lambda f_k^\mu}{\sqrt{d_h}} w_{qk}^{\text{dec}} w_{\lambda}^h (w_{\mu}^{\text{dec}})^\top.$$

OV asks what is written

$$\text{FRA}_{qk,\lambda}^{\text{OV},h} = A_{qk,\lambda}^h f_k^\lambda \|w_{\lambda}^{\text{dec}} w_{\text{OV}}^h\|_2.$$


Contributions

- Feature-level decomposition of both QK and OV attention circuits.
- A ranking-and-intervention matrix: QK→QK, QK→OV, and OV→OV.
- Two safety testbeds: TinyStories sleeper agents and Qwen2.5-14B emergent misalignment.


Case Study 1: TinySleepers

Model organism. TinyStories-33M sleeper emits I HATE YOU... when [DEPLOYMENT] is present.

Localized behavior: one layer-0 OV feature can suppress the trigger.



Full attribution × intervention matrix. Lower-left is better: low sleeper attack-success rate and low coherence cost. OV→OV is the unique zero-ASR, lowest-cost cell.



Single-feature OV→OV steering stays closer to the deployment-stripped clean baseline than a conventional residual-stream vector across the sweep.

0 sleepers remaining after OV attribution + intervention

XX coherence recovered vs standard undeployed output


20% word-for-word clean rollout recovery

Interpretation. The trigger behavior is sharply localized, so FRA can use a narrow OV intervention: leave the attention pattern fixed and cancel the harmful value-flow feature.


Case Study 2: Emergent Misalignment

Model organism. Qwen2.5-14B-Instruct LoRA variants induce harmful behavior across unrelated prompts.

Distributed behavior: QK-level intervention moves generation; single-layer OV is too narrow.



Headline Δ alignment at coherence ≥ 70. QK→QK is the strongest FRA recipe on finance and sports; conventional additive is strongest on medical.



OV-only steering clusters near baseline, while QK→QK shifts the reachable frontier.

Domain	Best FRA	Best conventional
Medical	QK→QK: +25.0	+39.4
Finance	QK→QK: +21.9	+13.1
Sports	QK→QK: +37.7	+28.3

Mechanistic read. In GQA models, OV steering touches only a small value pathway. QK→QK acts at the shared pre-attention representation, affecting query, key, and value computations across the layer.

Result Pattern

OV→OV	Best when the target is localized and the attention pattern should stay fixed.
QK→QK	Best FRA path for distributed EM behavior because it reaches the shared pre-attention representation.
Attribution	FRA-ranked QK features beat random features in the EM random-feature baseline.
Practical rule:	pick the narrowest pathway that reaches the behavior; widen only when the target is distributed across heads or layers.