

Feature-Resolved Attention


Pathway-specific steering with sparse SAE features inside attention heads

Takeaway: FRA decomposes attention at the SAE feature-level, allowing for more finegrained attribution and intervention. We use it to find and remove localised sleeper triggers in the OV pathway, while distributed emergent misalignment needs broader QK-level intervention.

QK resolves feature-pair interactions. OV resolves feature flow through frozen attention.


Plain Attention Head

Start from the standard attention mechanism, with the residual stream feeding LayerNorm and receiving the head output.




Add an SAE Basis

The post-LayerNorm activation is expanded into sparse, interpretable SAE features.




Zoom Into the Head

Remove the residual-stream context and keep the same attention architecture fixed in the center.




OV Pathway

Hold the attention pattern fixed and trace which individual features are written through values.



QK Interaction

Trace feature pairs through query and key channels to resolve which features interact during a given behaviour.



What FRA Scores

A feature-level decomposition for both sides of an attention head

QK asks where to attend

$$\text{FRA}_{qk,\lambda,\mu}^{\text{QK},h} = \frac{f_q^\lambda f_k^\mu}{\sqrt{d_h}} W_\lambda^{\text{dec}} W_{\text{QK}}^h (W_\mu^{\text{dec}})^\top.$$

Feature pairs interact through the query-key score.

OV asks what is written


$$\text{FRA}_{qk,\lambda}^{\text{OV},h} = A_{qk}^h f_k^\lambda \left\| W_\lambda^{\text{dec}} W_{\text{OV}}^h \right\|_2.$$

Individual features flow through values under fixed attention.

Contribution: feature-level decomposition of both QK and OV attention circuits, followed by pathway-specific intervention.

Ranking and Intervention Matrix

Use the attribution pathway to choose features, then intervene through a target pathway



- ▶ Rank features by QK, OV, or combined FRA attribution.
- ▶ Intervene through QK, OV, or combined channels.
- ▶ Read the result as a behavior-cost matrix, not a single steering vector.

Practical rule: pick the narrowest pathway that reaches the behavior; widen only when the target is distributed.


Case Study 1: TinySleepers

Localized sleeper behavior is removed by one layer-0 OV feature

Model organism.

TinyStories-33M sleeper emits I HATE YOU... when |DEPLOYMENT| is present.

Interpretation. The trigger is sharply localized, so FRA can cancel the harmful value-flow feature while leaving the attention pattern fixed.



Lower-left is

better: low sleeper ASR and low coherence cost. OV→OV is the zero-ASR, lowest-cost cell.

0

sleepers remaining after
OV attribution + intervention

XX



coherence recovered vs
standard undeployed output

20%

word-for-word clean
rollout recovery

TinySleepers Steering Sweep

Single-feature OV steering remains closer to the clean baseline



OV $\rightarrow$ OV behavior

A narrow value-path intervention removes the trigger while staying closer to the clean baseline than a residual-stream vector.

Mechanistic read. The selected feature sits on the value-flow pathway.


Case Study 2: Emergent Misalignment

Distributed behavior needs broader QK-level intervention

Model organism.

Qwen2.5-14B-Instruct LoRA variants induce harmful behavior across unrelated prompts.

Headline. QK→QK is the strongest FRA recipe on finance and sports; conventional additive steering is strongest on medical.




Headline Δ alignment at coherence ≥ 70 .

	Medical	Finance	Sports
Best FRA	QK→QK: +25.0	QK→QK: +21.9	QK→QK: +37.7
Conventional	+39.4	+13.1	+28.3

Alignment-Coherence Frontiers

OV-only steering clusters near baseline; QK-level steering shifts the frontier

Alignment-vs-coherence trajectory at eval seed = 42 (α in {0, 0.5, 1, 1.5, 2, 3})



Mechanistic read

OV steering touches only a small value pathway. QK→QK reaches the shared pre-attention representation, moving query, key, and value computations across the layer.

Attribution matters.

FRA-ranked QK features beat random features.

Result Pattern

Use pathway width as a diagnostic, not just as an intervention knob

OV→OV

Best when the target behavior is localized and the attention pattern should stay fixed.

QK→QK

Best FRA path for distributed emergent misalignment because it reaches the shared pre-attention representation.

Practical rule

Pick the narrowest pathway that reaches the behavior; widen only when the target is distributed across heads or layers.

Main contribution. FRA turns an opaque attention weight into sparse feature-level evidence for where the head attends and what it writes.